Evidence-Calibrated Self-Correcting RAG — Guide Results

Measured Linux CPU pilot

Built: bounded hybrid retrieval, extractive candidate selection, project-trained answer-bearing-evidence proxy, one retrieval-changing correction, exact-match calibration, citations, and abstention.

method	em	f1	coverage	selective_risk	recall_at_5	median_latency_ms	p95_latency_ms	mean_retrieval_calls	mean_generation_calls
B0_closed_bo	0.0	0.02	0.0	nan	0.0	0.0	0.0	0.0	0.0
B1_bm25	0.02	0.148	1.0	0.98	0.587	1.471	1.802	1.0	1.0
B2_hybrid	0.013	0.126	1.0	0.987	0.5	1.589	2.165	1.0	1.0
B3_corrected	0.013	0.13	1.0	0.987	0.513	3.471	4.224	1.8	1.8
B4_full	0.0	0.02	0.0	nan	0.513	3.548	4.408	1.8	1.8

Calibration: ECE 0.195; Brier 0.049; AURC 0.993.

Important limitations: rule-based extractive generator, support proxy rather than entailment, bounded SQuAD corpus, no human labels, Linux-only execution. Full-system zero coverage is a negative result, not a zero-risk guarantee.

Inspect results/dashboard.html, plots/, tables/, example_outputs.json, and per_question_predictions.jsonl.